

# House Prices Regression Task

Matthew Loh

## Table of contents

Data Loading and Exploratory Data Analysis . . . . .	1
Data Preprocessing . . . . .	9
Model Training and Evaluation . . . . .	9
Model Performance Visualization . . . . .	13
Feature Importance . . . . .	18
Hyperparameter Tuning . . . . .	21
Final Model Evaluation . . . . .	25
Predictions on Test Data . . . . .	31
Conclusion . . . . .	32

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score, accuracy_score
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.ensemble import RandomForestRegressor
import lightgbm as lgb
import xgboost as xgb
import catboost as cbt

np.random.seed(42)
```

## Data Loading and Exploratory Data Analysis

```
# Load the data
train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')

# Display basic information about the dataset
print(train.info())

# Display summary statistics
print(train.describe())

# Plot distribution of target variable
plt.figure(figsize=(10, 6))
sns.histplot(train['SalePrice'], kde=True)
plt.title('Distribution of Sale Prices')
plt.show()

# Identify numeric columns
numeric_columns = train.select_dtypes(include=[np.number]).columns

# Correlation matrix of numerical features
corr_matrix = train[numeric_columns].corr()
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, cmap='coolwarm', annot=False)
plt.title('Correlation Matrix of Numerical Features')
plt.show()

# Top 10 features correlated with SalePrice
top_corr = corr_matrix['SalePrice'].sort_values(ascending=False).head(11)
plt.figure(figsize=(10, 6))
sns.barplot(x=top_corr.index[1:], y=top_corr.values[1:])
plt.title('Top 10 Features Correlated with SalePrice')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

# Scatter plot of top correlated feature vs SalePrice
top_feature = top_corr.index[1]
plt.figure(figsize=(10, 6))
sns.scatterplot(x=train[top_feature], y=train['SalePrice'])
```

```
plt.title(f'{top_feature} vs SalePrice')
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea               1460 non-null   int64
5   Street               1460 non-null   object
6   Alley                91 non-null     object
7   LotShape              1460 non-null   object
8   LandContour           1460 non-null   object
9   Utilities             1460 non-null   object
10  LotConfig             1460 non-null   object
11  LandSlope             1460 non-null   object
12  Neighborhood          1460 non-null   object
13  Condition1            1460 non-null   object
14  Condition2            1460 non-null   object
15  BldgType              1460 non-null   object
16  HouseStyle            1460 non-null   object
17  OverallQual           1460 non-null   int64
18  OverallCond           1460 non-null   int64
19  YearBuilt             1460 non-null   int64
20  YearRemodAdd          1460 non-null   int64
21  RoofStyle             1460 non-null   object
22  RoofMatl              1460 non-null   object
23  Exterior1st           1460 non-null   object
24  Exterior2nd           1460 non-null   object
25  MasVnrType            588 non-null     object
26  MasVnrArea            1452 non-null   float64
27  ExterQual             1460 non-null   object
28  ExterCond             1460 non-null   object
29  Foundation            1460 non-null   object
30  BsmtQual              1423 non-null   object
31  BsmtCond              1423 non-null   object
32  BsmtExposure          1422 non-null   object
33  BsmtFinType1          1423 non-null   object
```

34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	7	non-null	object
73	Fence	281	non-null	object
74	MiscFeature	54	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64

```

77  YrSold          1460 non-null   int64
78  SaleType        1460 non-null   object
79  SaleCondition    1460 non-null   object
80  SalePrice        1460 non-null   int64

```

dtypes: float64(3), int64(35), object(43)

memory usage: 924.0+ KB

None

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	\
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	
std	421.610009	42.300571	24.284752	9981.264932	1.382997	
min	1.000000	20.000000	21.000000	1300.000000	1.000000	
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	


	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	...	\
count	1460.000000	1460.000000	1460.000000	1452.000000	1460.000000	...	
mean	5.575342	1971.267808	1984.865753	103.685262	443.639726	...	
std	1.112799	30.202904	20.645407	181.066207	456.098091	...	
min	1.000000	1872.000000	1950.000000	0.000000	0.000000	...	
25%	5.000000	1954.000000	1967.000000	0.000000	0.000000	...	
50%	5.000000	1973.000000	1994.000000	0.000000	383.500000	...	
75%	6.000000	2000.000000	2004.000000	166.000000	712.250000	...	
max	9.000000	2010.000000	2010.000000	1600.000000	5644.000000	...	


	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	94.244521	46.660274	21.954110	3.409589	15.060959	
std	125.338794	66.256028	61.119149	29.317331	55.757415	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	25.000000	0.000000	0.000000	0.000000	
75%	168.000000	68.000000	0.000000	0.000000	0.000000	
max	857.000000	547.000000	552.000000	508.000000	480.000000	



	PoolArea	MiscVal	MoSold	YrSold	SalePrice
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	2.758904	43.489041	6.321918	2007.815753	180921.195890
std	40.177307	496.123024	2.703626	1.328095	79442.502883
min	0.000000	0.000000	1.000000	2006.000000	34900.000000
25%	0.000000	0.000000	5.000000	2007.000000	129975.000000

50%	0.000000	0.000000	6.000000	2008.000000	163000.000000
75%	0.000000	0.000000	8.000000	2009.000000	214000.000000
max	738.000000	15500.000000	12.000000	2010.000000	755000.000000

[8 rows x 38 columns]









## Data Preprocessing

```
def preprocess_data(df):
    # Handle missing values
    for col in df.columns:
        if df[col].dtype != 'object':
            df[col] = df[col].fillna(df[col].median())
        else:
            df[col] = df[col].fillna(df[col].mode()[0])

    # Encode categorical variables
    le = LabelEncoder()
    for col in df.select_dtypes(include=['object']).columns:
        df[col] = le.fit_transform(df[col].astype(str))

    return df

# Preprocess train and test data
X = preprocess_data(train.drop('SalePrice', axis=1))
y = np.log1p(train['SalePrice']) # Log transform the target variable
test_processed = preprocess_data(test)

# Apply StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
test_processed_scaled = scaler.transform(test_processed)

print("Processed data shape:", X.shape)
```

Processed data shape: (1460, 80)

## Model Training and Evaluation

```
def train_and_evaluate(model, X, y, test_data, model_name):
    X_train, X_val, y_train, y_val = train_test_split(
        X, y, test_size=0.2, random_state=42)

    model.fit(X_train, y_train)
```

```

train_pred = model.predict(X_train)
val_pred = model.predict(X_val)

train_mse = mean_squared_error(y_train, train_pred)
train_rmse = np.sqrt(train_mse)
train_mae = mean_absolute_error(y_train, train_pred)
train_r2 = r2_score(y_train, train_pred)

val_mse = mean_squared_error(y_val, val_pred)
val_rmse = np.sqrt(val_mse)
val_mae = mean_absolute_error(y_val, val_pred)
val_r2 = r2_score(y_val, val_pred)

print(f"{model_name} Results:")
print(f"Train RMSE: {train_rmse:.4f}")
print(f"Train MAE: {train_mae:.4f}")
print(f"Train R2 Score: {train_r2:.4f}")
print(f"Validation RMSE: {val_rmse:.4f}")
print(f"Validation MAE: {val_mae:.4f}")
print(f"Validation R2 Score: {val_r2:.4f}")
print("\n")

return model, (y_train, train_pred, y_val, val_pred)

# Linear models
linear_models = {
    "Linear Regression": LinearRegression(),
    "Ridge": Ridge(),
    "Lasso": Lasso(),
    "ElasticNet": ElasticNet()
}

linear_results = {}
for name, model in linear_models.items():
    linear_results[name] = train_and_evaluate(
        model, X_scaled, y, test_processed_scaled, name)

# Advanced models
rf_model = RandomForestRegressor(random_state=42)
rf_trained, rf_results = train_and_evaluate(
    rf_model, X, y, test_processed, "Random Forest")

```

```

lgb_model = lgb.LGBMRegressor(random_state=42)
lgb_trained, lgb_results = train_and_evaluate(
    lgb_model, X, y, test_processed, "LightGBM")

xgb_model = xgb.XGBRegressor(random_state=42)
xgb_trained, xgb_results = train_and_evaluate(
    xgb_model, X, y, test_processed, "XGBoost")

cbt_model = cbt.CatBoostRegressor(random_state=42, verbose=False)
cbt_trained, cbt_results = train_and_evaluate(
    cbt_model, X, y, test_processed, "CatBoost")

```

#### Linear Regression Results:

Train RMSE: 0.1306  
 Train MAE: 0.0894  
 Train R2 Score: 0.8882  
 Validation RMSE: 0.1553  
 Validation MAE: 0.1061  
 Validation R2 Score: 0.8708

#### Ridge Results:

Train RMSE: 0.1306  
 Train MAE: 0.0894  
 Train R2 Score: 0.8881  
 Validation RMSE: 0.1553  
 Validation MAE: 0.1061  
 Validation R2 Score: 0.8708

#### Lasso Results:

Train RMSE: 0.3904  
 Train MAE: 0.3034  
 Train R2 Score: 0.0000  
 Validation RMSE: 0.4332  
 Validation MAE: 0.3371  
 Validation R2 Score: -0.0058

#### ElasticNet Results:

Train RMSE: 0.3904

Train MAE: 0.3034  
Train R2 Score: 0.0000  
Validation RMSE: 0.4332  
Validation MAE: 0.3371  
Validation R2 Score: -0.0058

Random Forest Results:

Train RMSE: 0.0535  
Train MAE: 0.0361  
Train R2 Score: 0.9812  
Validation RMSE: 0.1458  
Validation MAE: 0.0986  
Validation R2 Score: 0.8861

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.0014.  
You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 3364

[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 73

[LightGBM] [Info] Start training from score 12.030658

LightGBM Results:

Train RMSE: 0.0420  
Train MAE: 0.0256  
Train R2 Score: 0.9884  
Validation RMSE: 0.1413  
Validation MAE: 0.0934  
Validation R2 Score: 0.8930

XGBoost Results:

Train RMSE: 0.0050  
Train MAE: 0.0035  
Train R2 Score: 0.9998  
Validation RMSE: 0.1514  
Validation MAE: 0.1022  
Validation R2 Score: 0.8772



CatBoost Results:


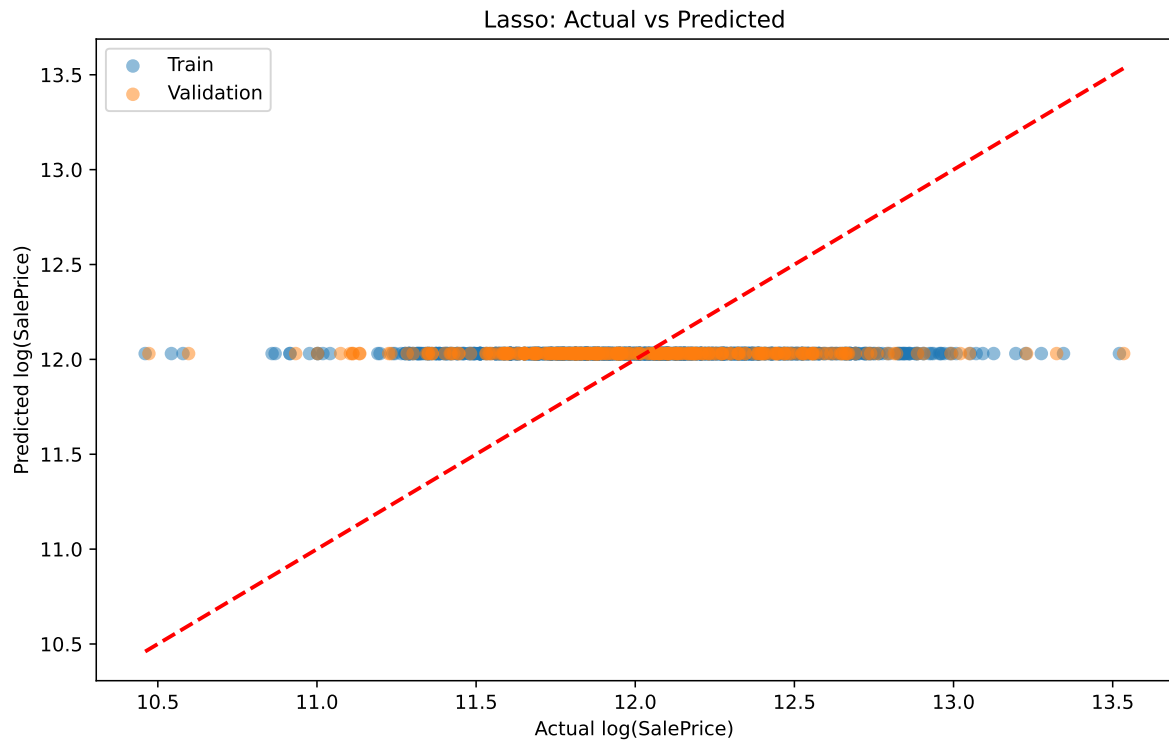
Train RMSE: 0.0322  
Train MAE: 0.0246  
Train R2 Score: 0.9932

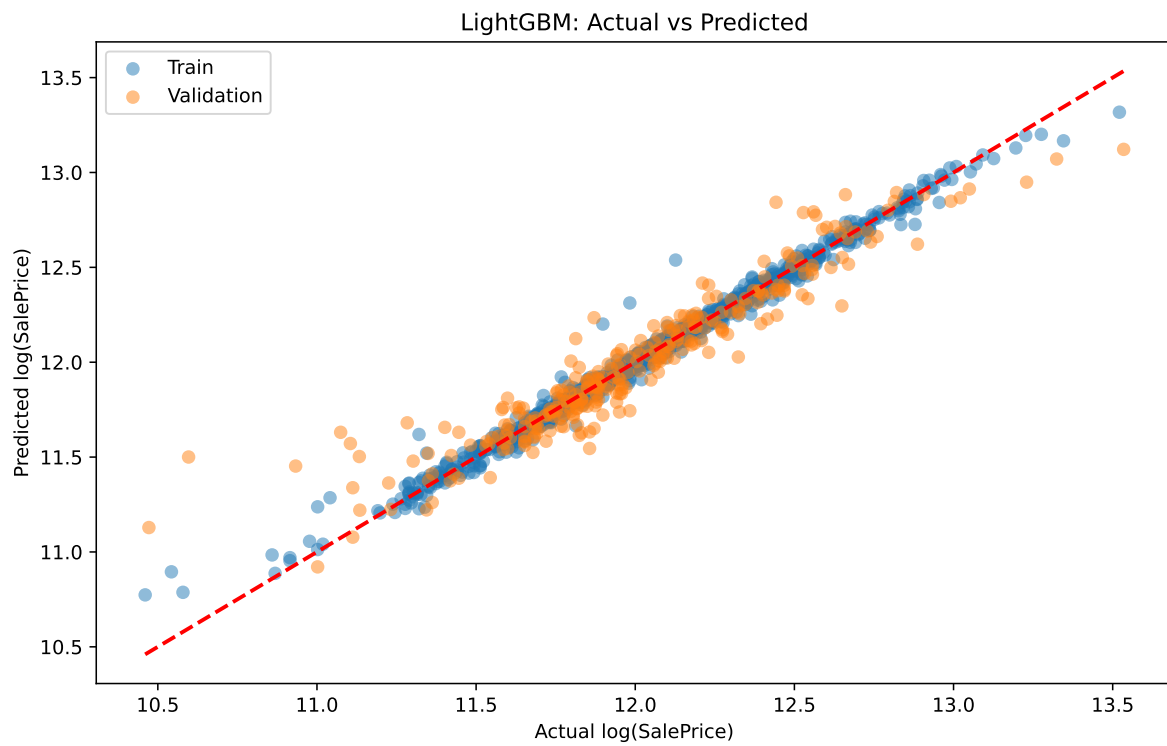

Validation RMSE: 0.1329  
Validation MAE: 0.0866  
Validation R2 Score: 0.9053

## Model Performance Visualization



```
def plot_actual_vs_predicted(results, model_name):  
    y_train, train_pred, y_val, val_pred = results  
    plt.figure(figsize=(10, 6))  
    plt.scatter(y_train, train_pred, alpha=0.5, label='Train')  
    plt.scatter(y_val, val_pred, alpha=0.5, label='Validation')  
    plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', lw=2)  
    plt.xlabel('Actual log(SalePrice)')  
    plt.ylabel('Predicted log(SalePrice)')  
    plt.title(f'{model_name}: Actual vs Predicted')  
    plt.legend()  
    plt.show()  
  
# Plot for each model  
for name, (model, results) in linear_results.items():  
    plot_actual_vs_predicted(results, name)  
  
plot_actual_vs_predicted(rf_results, "Random Forest")  
plot_actual_vs_predicted(lgb_results, "LightGBM")  
plot_actual_vs_predicted(xgb_results, "XGBoost")  
plot_actual_vs_predicted(cbt_results, "CatBoost")
```











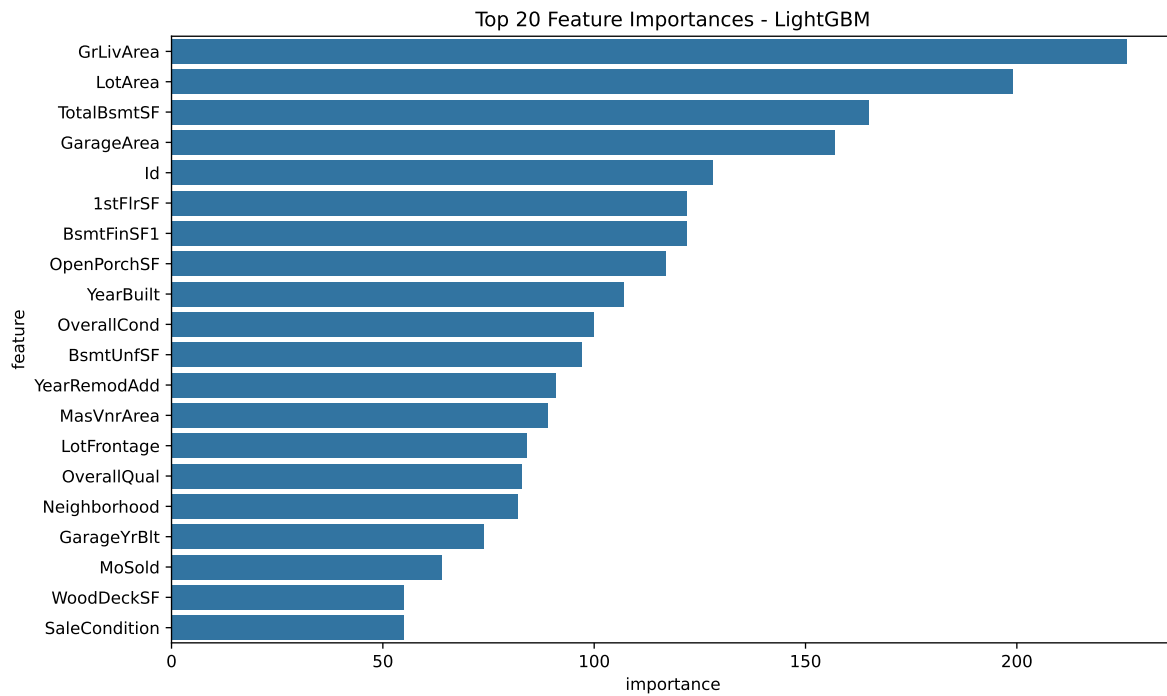

## Feature Importance

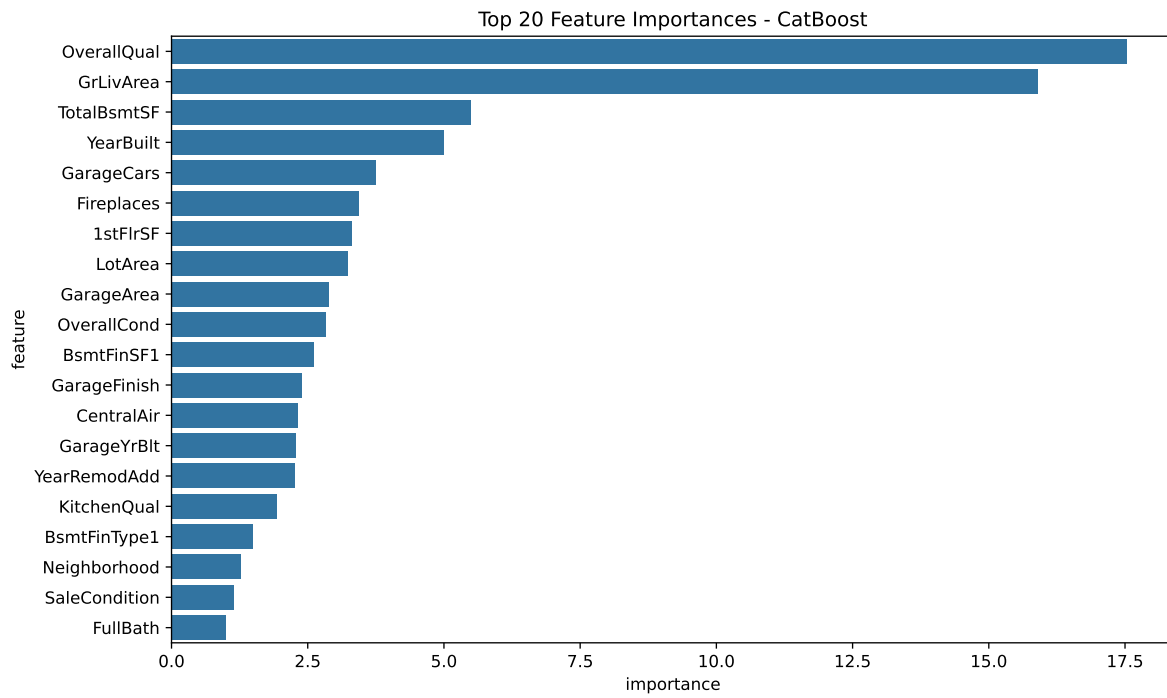

```
def plot_feature_importance(model, X, model_name):
    if hasattr(model, 'feature_importances_'):
        importances = model.feature_importances_
    elif hasattr(model, 'feature_importance'):
        importances = model.feature_importance()
    else:
        print(f"Feature importance not available for {model_name}")
        return

    feature_imp = pd.DataFrame(
        {'feature': X.columns, 'importance': importances})
    feature_imp = feature_imp.sort_values(
        'importance', ascending=False).head(20)

    plt.figure(figsize=(10, 6))
    sns.barplot(x='importance', y='feature', data=feature_imp)
    plt.title(f'Top 20 Feature Importances - {model_name}')
    plt.tight_layout()
    plt.show()

plot_feature_importance(rf_trained, X, "Random Forest")
plot_feature_importance(lgb_trained, X, "LightGBM")
plot_feature_importance(xgb_trained, X, "XGBoost")
plot_feature_importance(cbt_trained, X, "CatBoost")
```





## Hyperparameter Tuning

```
def tune_hyperparameters(model, param_grid, X, y, model_name):
    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grid,
        cv=5,
        scoring="neg_mean_squared_error",
        verbose=1,
        n_jobs=-1,
    )
    grid_search.fit(X, y)

    print(f"Best parameters for {model_name}:")
    print(grid_search.best_params_)
    print(f"Best RMSE: {np.sqrt(-grid_search.best_score_):.4f}")
    print("\n")

    return grid_search.best_estimator_

# Random Forest hyperparameter tuning
rf_param_grid = {
    "n_estimators": [100, 200],
    "max_depth": [None, 10],
    "min_samples_split": [2, 5],
}
rf_tuned = tune_hyperparameters(
    RandomForestRegressor(random_state=42), rf_param_grid, X, y, "Random Forest"
)

# LightGBM hyperparameter tuning
lgb_param_grid = {
    "num_leaves": [31, 127],
    "learning_rate": [0.01, 0.1],
    "n_estimators": [100, 200],
}
lgb_tuned = tune_hyperparameters(
    lgb.LGBMRegressor(random_state=42), lgb_param_grid, X, y, "LightGBM"
)

# XGBoost hyperparameter tuning
```

```

xgb_param_grid = {
    "max_depth": [3, 6],
    "learning_rate": [0.01, 0.1],
    "n_estimators": [100, 200],
}
xgb_tuned = tune_hyperparameters(
    xgb.XGBRegressor(random_state=42), xgb_param_grid, X, y, "XGBoost"
)

# CatBoost hyperparameter tuning
cbt_param_grid = {
    "depth": [6, 8],
    "learning_rate": [0.01, 0.1],
    "iterations": [100, 200],
}
cbt_tuned = tune_hyperparameters(
    cbt.CatBoostRegressor(random_state=42, verbose=False),
    cbt_param_grid,
    X,
    y,
    "CatBoost",
)

```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

Best parameters for Random Forest:

```
{'max_depth': None, 'min_samples_split': 5, 'n_estimators': 200}
```

Best RMSE: 0.1423

Fitting 5 folds for each of 8 candidates, totalling 40 fits

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.0026s

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 3625

[LightGBM] [Info] Number of data points in the train set: 1460, number of used features: 74

[LightGBM] [Info] Start training from score 12.024057

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf



[illegible]



```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
Best parameters for LightGBM:
{'learning_rate': 0.1, 'n_estimators': 100, 'num_leaves': 127}
Best RMSE: 0.1338
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

Best parameters for XGBoost:

```
{'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 200}
```

Best RMSE: 0.1271

Fitting 5 folds for each of 8 candidates, totalling 40 fits

Best parameters for CatBoost:

```
{'depth': 6, 'iterations': 200, 'learning_rate': 0.1}
```

Best RMSE: 0.1248

## Final Model Evaluation

```
print("Final Model Evaluation:")
rf_final, rf_final_results = train_and_evaluate(
    rf_tuned, X, y, test_processed, "Random Forest (Tuned)"
)
lgb_final, lgb_final_results = train_and_evaluate(
    lgb_tuned, X, y, test_processed, "LightGBM (Tuned)"
)
xgb_final, xgb_final_results = train_and_evaluate(
    xgb_tuned, X, y, test_processed, "XGBoost (Tuned)"
)
cbt_final, cbt_final_results = train_and_evaluate(
    cbt_tuned, X, y, test_processed, "CatBoost (Tuned)"
)

# Plot final model performances
```

```
plot_actual_vs_predicted(rf_final_results, "Random Forest (Tuned)")
plot_actual_vs_predicted(lgb_final_results, "LightGBM (Tuned)")
plot_actual_vs_predicted(xgb_final_results, "XGBoost (Tuned)")
plot_actual_vs_predicted(cbt_final_results, "CatBoost (Tuned)")
```

Final Model Evaluation:

Random Forest (Tuned) Results:

Train RMSE: 0.0597

Train MAE: 0.0395

Train R2 Score: 0.9766

Validation RMSE: 0.1463

Validation MAE: 0.0984

Validation R2 Score: 0.8853

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.0006s.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 3364

[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 73

[LightGBM] [Info] Start training from score 12.030658

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[illegible]



XGBoost (Tuned) Results:

Train RMSE: 0.0622

Train MAE: 0.0456

Train R2 Score: 0.9746

Validation RMSE: 0.1372

Validation MAE: 0.0919

Validation R2 Score: 0.8991

CatBoost (Tuned) Results:

Train RMSE: 0.0606

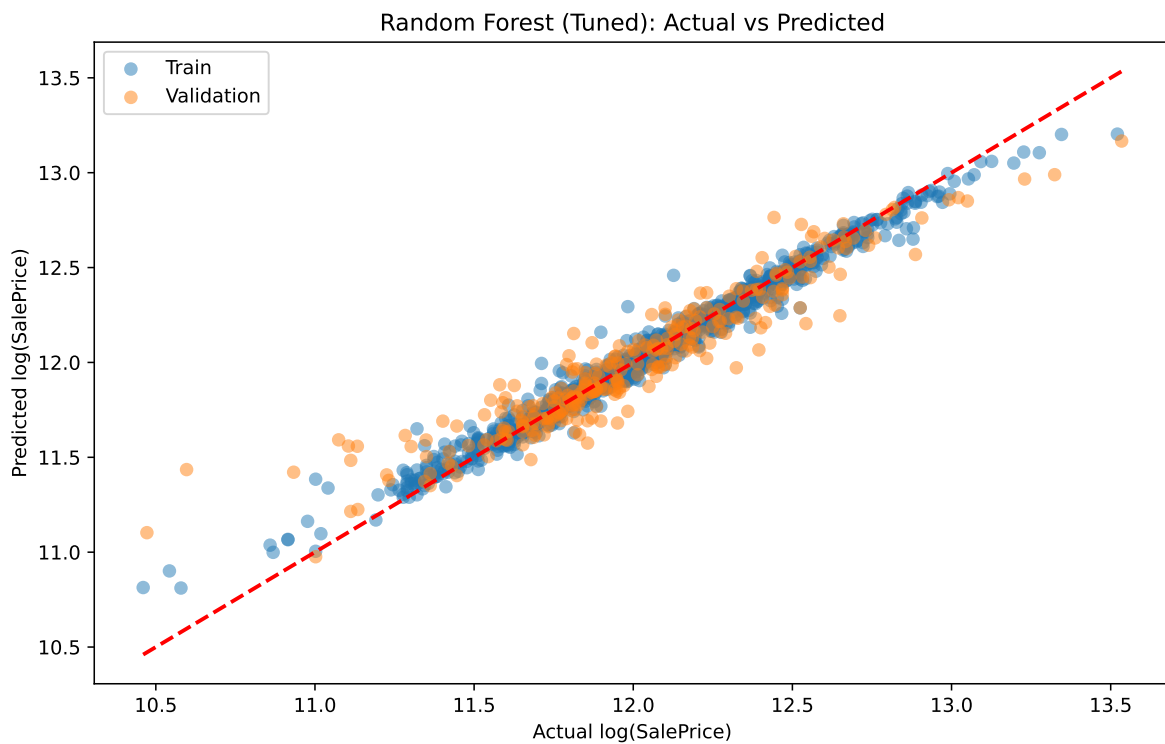
Train MAE: 0.0450

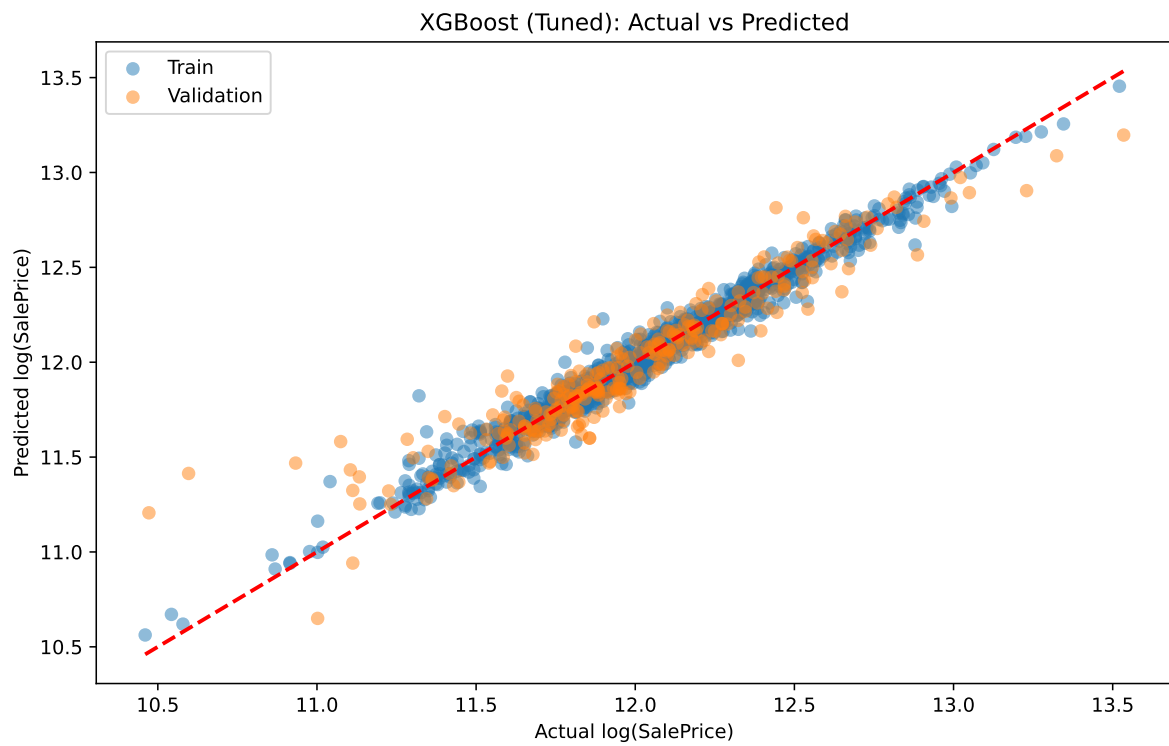
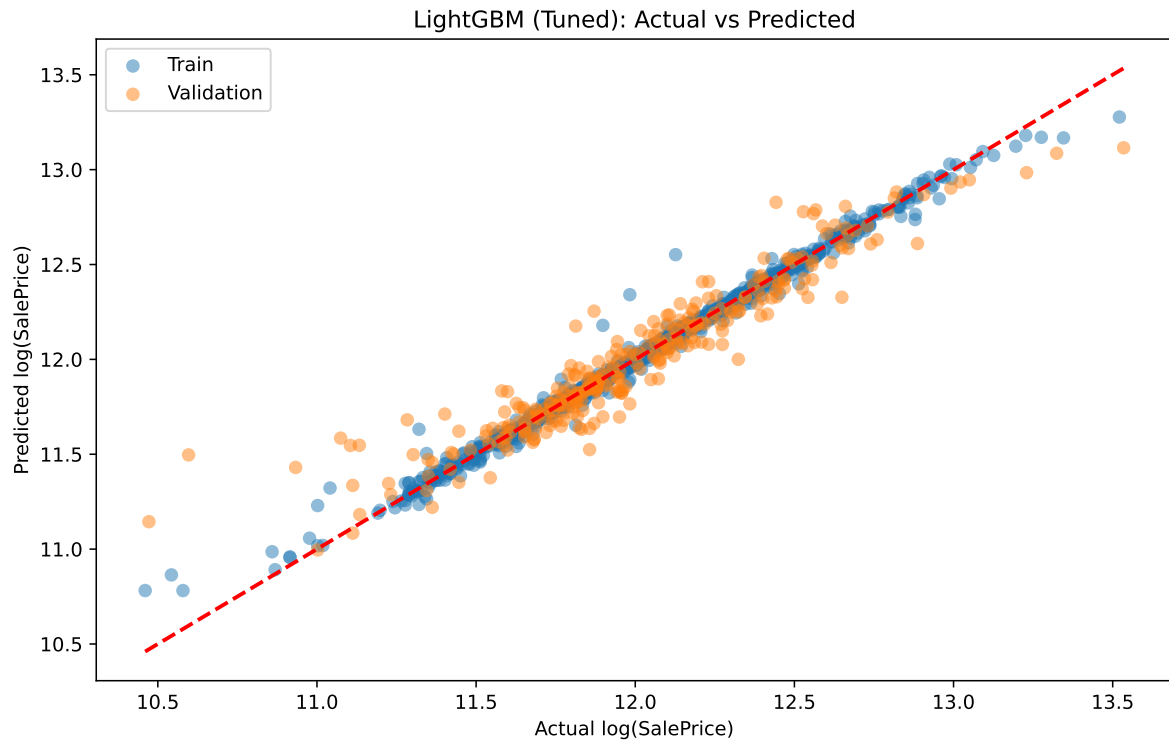
Train R2 Score: 0.9759

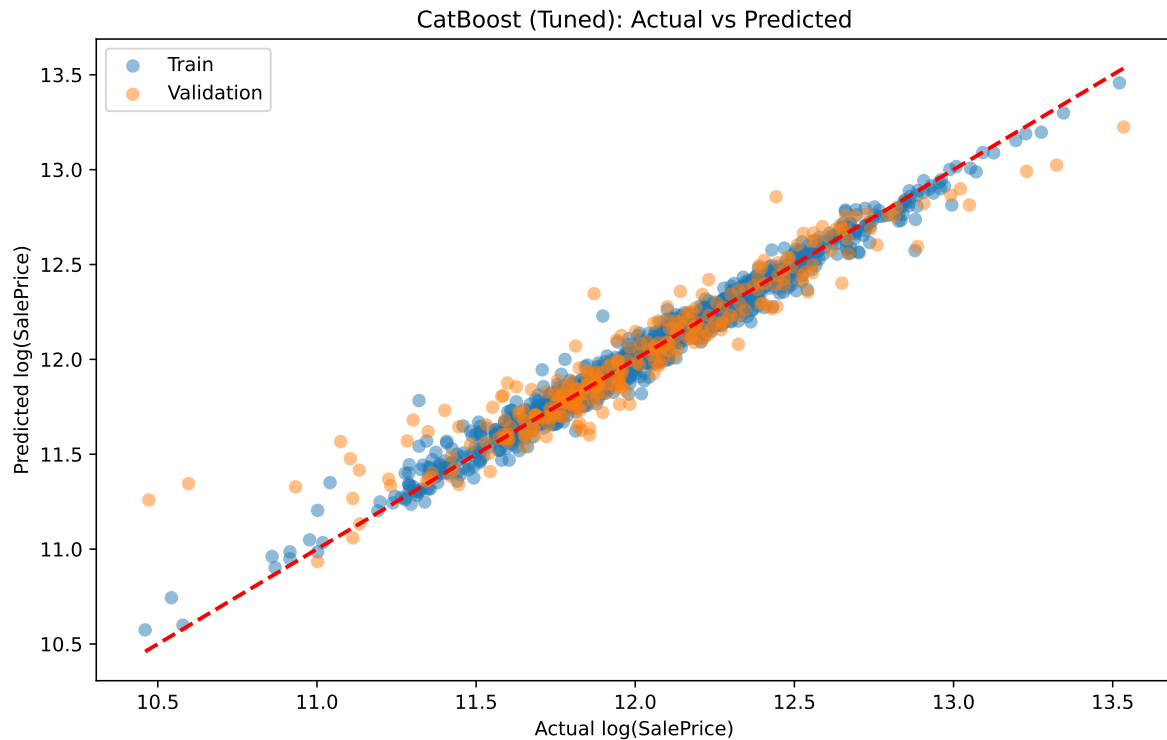
Validation RMSE: 0.1334

Validation MAE: 0.0871

Validation R2 Score: 0.9047







## Predictions on Test Data

```
def make_predictions(model, test_data):  
    predictions = model.predict(test_data)  
    return np.expml(predictions) # Inverse log transform  
  
rf_predictions = make_predictions(rf_final, test_processed)  
lgb_predictions = make_predictions(lgb_final, test_processed)  
xgb_predictions = make_predictions(xgb_final, test_processed)  
cbr_predictions = make_predictions(cbr_final, test_processed)  
  
# Ensemble predictions (simple average)  
ensemble_predictions = (rf_predictions + lgb_predictions +  
                        xgb_predictions + cbr_predictions) / 4  
  
# Plot distribution of predictions  
plt.figure(figsize=(12, 6))  
sns.kdeplot(rf_predictions, label='Random Forest')
```


```

sns.kdeplot(lgb_predictions, label='LightGBM')
sns.kdeplot(xgb_predictions, label='XGBoost')
sns.kdeplot(cbt_predictions, label='CatBoost')
sns.kdeplot(ensemble_predictions, label='Ensemble')
plt.xlabel('Predicted SalePrice')
plt.ylabel('Density')
plt.title('Distribution of Predicted Sale Prices')
plt.legend()
plt.show()

# Make predictions on the test data
test_predictions = make_predictions(rf_final, test_processed)

# Create submission file
submission = pd.DataFrame(
    {'Id': test['Id'], 'SalePrice': test_predictions})
submission.to_csv('submission.csv', index=False)

```



## Conclusion

This notebook implements a comprehensive approach to the House Prices regression task, including:



1. Exploratory Data Analysis (EDA) to understand the dataset
2. Data preprocessing, including handling missing values and encoding categorical variables
3. Implementation of both basic (linear) and advanced (tree-based) regression models
4. Visualization of model performance and feature importance
5. Hyperparameter tuning to optimize model performance
6. Final model evaluation and ensemble prediction

Key observations: 1. The log transformation of the target variable (SalePrice) helped to handle its skewed distribution. 2. Advanced models (Random Forest, LightGBM, XGBoost, CatBoost) generally outperformed linear models. 3. Feature importance analysis revealed key predictors of house prices, which align with domain knowledge. 4. Hyperparameter tuning improved the performance of all models. 5. The ensemble of tuned models provides a robust final prediction.

Areas for further improvement: 1. More extensive feature engineering, such as creating interaction terms or domain-specific features. 2. Experimenting with more advanced ensemble methods, such as stacking. 3. Deeper analysis of residuals to identify patterns in prediction errors and potential outliers. 4. Consideration of model interpretability for stakeholder communication.